

CSE 842: NATURAL LANGUAGE PROCESSING HOMEWORK 3 ASSIGNMENT

Due Date: November 20, 2024 at 11:59 PM

OVERVIEW

For this assignment, I want everyone to become familiar with working with BERT. I have linked to tutorials taken from HuggingFace below. The [HuggingFace Tutorials](#) are linked to Google Collab, so you can use Collab's GPUs to run your code. I've specified [BERT](#) throughout the assignment. However, you may use any BERT variant you prefer. [RoBERTa](#) works really well for NLP tasks; [DistilBERT](#) is a smaller, faster version of BERT. You can use PyTorch or Tensorflow.

PROBLEM 1: FINE-TUNE A PRETRAINED MODEL

Description: For this problem, you are to walk through the HuggingFace tutorial for fine tuning: [Fine Tuning a Pretrained Model](#). The tutorial has sections for both PyTorch & TensorFlow. This is a simple fine-tuning approach that will not require you to alter the transformer layers of the pretrained model.

What to submit:

1. Code: Completed Python file of the tutorial
2. Written Report: A brief description including:
 - a. the package you used (PyTorch, TensorFlow)
 - b. which model you used (The tutorial uses *bert-base-cased*, you should replace this with a BERT model of your choosing. I suggest running the tutorial as is first.)
 - c. brief description of model (`model.summary()`)
 - d. your results
 - e. something new or interesting that you learned during implementation.

PROBLEM 2: OUTPERFORM BERT

Description: For this problem, you can choose to solve any NLP task we've discussed this semester. For example: PoS tagging, NER, sentiment classification, topic classification, QA, MT, summarization, SRL, etc. Your task is to develop a model (or modeling pipeline) and relevant features to try to outperform the current best BERT-related-model performance on your chosen task. *Your limitation is: you can't use BERT, transformers, encoders, or neural networks of any kind.* If you want to use word embeddings, you can use [word2vec](#), [GloVe](#), or [ELMO](#).

What to submit:

1. Code: Python file of your model.
2. Written Report: Provide instructions on how to run your code as in previous homework assignments. Provide details on which task, features, model, etc. you chose and **why** you chose them. Include the following:
 - a. Your chosen NLP task and why you wanted to study it.

- b. The dataset you are using and why you chose it for this task.
(For example, if you want to work with NER tagging, you probably don't want to choose a Twitter dataset labeled for sentiment.)
- c. The model(s) you used. You may use any non-NN model: logistic regression, SVM, HMM, CRF, etc. You do not need to implement these from scratch. You can use scikit-learn or any other packages you need. However, do not use a pre-existing paper's entire code base including model & features.
- d. The features you tested. For (c) & (d), I would suggest looking up the previous (not NN) state-of-the-art model for your task and using that to initialize your ideas. Link to that paper if you use this approach.
- e. Link to the most recent, best-performing BERT (or related) model for this task. Include a table summary of their results (from their paper).
- f. Summary table of your results (precision, recall, F1, accuracy).
- g. Brief discussion of:
 - i. where you think your modeling approach and features worked well
 - ii. where you might have been able to improve
 - iii. why you think your model did or did not outperform the BERT model of (e).
- h. *Optional for 5 bonus points:* Add BERT contextual embeddings into your features and report/discuss your new results.
- i. *Optional for 5 bonus points:* Replace your model in (c) with a bidirectional LSTM. Report/discuss your new results.

PROBLEM 3: LAST 10 POINTS

Problem 1 is worth 10 points. Problem 2 is worth 80. To get the remaining 10 points, you must fine-tune BERT (or variant) to solve your task. You may want to limit your task selection to those that HuggingFace provides an *AutoModelFor-Task-* for you to use. This step may be challenging or time consuming depending on your prior NN experience. So I've provided the 2 bonus opportunities above (2h & 2i). If you complete everything, any excess bonus points can be applied to elsewhere in your grades.

SUBMISSION SUMMARY

Please submit your Python files and *one* written report in **PDF** format.
Please zip your files together and include your last name in the zipped filename.